

The Bot Club

AI Governance Evidence Pack

APRA CPS 230 · EU AI Act · ISO/IEC 42001

REVIEW WINDOW	90d
GENERATED AT	2026-05-15 13:47 UTC
GENERATOR VERSION	v1.0.0
SOURCE SHA-256	853aa49c9ff4e449840e9d90075f21620ad53b7a307c01f0afd476c0e429aafe
MERKLE ROOT	642ad10aefb062d4c54d54287d1aef2827cdd5c98854be49fa111afe9c647cf9
BTC ANCHOR	pending

Contents

- 1. Executive Summary
- 2. Scope and Agent Inventory
- 3. Risk Taxonomy and Control Map
- 4. Statistical Sampling
- 5. Incident and Gap Log
- 6. Third-Party and Model Registry
- 7. Spend and Rate-Limit Attestation
- 8. Tool-Call Inventory and Approval Log
- 9. Signature and Anchor Proof

EXECUTIVE SUMMARY

This evidence pack covers the period **90d**, generated on 2026-05-15 13:47 UTC for **The Bot Club** under the **APRA CPS 230 · EU AI Act · ISO/IEC 42001** framework(s). The AgentGuard platform evaluated 4,932 agent actions across 18 AI agents. The audit log hash chain was verified (pass), and the source dataset is cryptographically anchored in this document.

METRIC	VALUE	NOTES
Total agent evaluations	4,932	Allow 96.0% · Block 2.9% · HITL 1.0%
AI agents in scope	18	Across 3 framework(s)
Controls assessed	73	Pass 73 · Partial 0 · Gap 0
Overall conformance	100.0%	Target: ≥85% for compliance certification
HITL escalations	50	Approved 36 · Rejected 14
Incidents (total/resolved)	35 / 30	Remediation rate: 85.7%
Total AI spend (window)	USD 29.6147	Per-agent breakdown in Section 8

No control gaps identified in this review window.

SCOPE AND AGENT INVENTORY

The following 18 AI agents were active during the review window. Criticality ratings drive the stringency of controls applied (critical and high agents receive enhanced logging and mandatory HITL gates).

Agent ID	Runtime	Policy	Owner	Criticality	Model
brand-content-1	OpenClaw	policy-brand-v2.1	lucy.chen	MEDIUM	openai gpt-4o
brand-content-2	OpenClaw	policy-brand-v2.1	lucy.chen	MEDIUM	openai gpt-4o-mini
sales-research-1	OpenClaw	policy-sales-v3.0	james.okafor	HIGH	openai gpt-4o
sales-research-2	OpenClaw	policy-sales-v3.0	james.okafor	HIGH	openai gpt-4o
sales-outreach-1	LangChain	policy-sales-outreach-v1.5	james.okafor	HIGH	anthropic claude-3.5-sonnet
ops-dispatch-1	OpenClaw	policy-ops-critical-v4.2	priya.sharma	CRITICAL	openai gpt-4o
ops-dispatch-2	OpenClaw	policy-ops-critical-v4.2	priya.sharma	CRITICAL	openai gpt-4o
ops-dispatch-3	CrewAI	policy-ops-standard-v2.0	priya.sharma	HIGH	anthropic claude-3.5-haiku
data-pipeline-1	OpenClaw	policy-data-v2.3	dev	HIGH	openai gpt-4o-mini
data-pipeline-2	OpenClaw	policy-data-v2.3	dev	MEDIUM	openai gpt-4o-mini
finance-reconcile-1	OpenClaw	policy-finance-critical-v5.1	accounts	CRITICAL	openai gpt-4o
finance-reconcile-2	OpenAI Assistants	policy-finance-critical-v5.1	accounts	CRITICAL	openai gpt-4o
customer-support-1	OpenClaw	policy-support-v3.2	support	HIGH	anthropic claude-3.5-haiku
customer-support-2	OpenClaw	policy-support-v3.2	support	HIGH	anthropic claude-3.5-haiku
customer-support-3	MCP	policy-support-lite-v1.0	support	MEDIUM	anthropic claude-3.5-haiku
infra-monitor-1	OpenClaw	policy-infra-v2.7	ops	HIGH	openai gpt-4o-mini
infra-monitor-2	AutoGen	policy-infra-v2.7	ops	MEDIUM	google gemini-1.5-flash
research-analyst-1	OpenClaw	policy-research-v1.8	research	MEDIUM	anthropic claude-3.5-sonnet

RISK TAXONOMY AND CONTROL MAP

The following table maps each control in the selected framework(s) to observed evidence. Results are determined by evaluating the AgentGuard audit log against each control's evidence query. Sampled events and implementation statements are provided for each control.

CPS230

Control ID	Clause	Requirement	Evidence / Implementation	Result	Reviewer
CPS230-GOV-01	22	The Board must approve and maintain an operational risk management framework that is commensurate with the nature, scale and complexity of the entity's operations, including AI-driven processes.	AgentGuard records every policy attachment, policy version change, and approval workflow for AI agent deployments. Policy lineage is immutable and board-approved policy IDs are captured in the agent configuration manifest. All 18 agents registered with board-approved policy versions.	✓ PASS	Chief Risk Officer
CPS230-GOV-02	23	Senior management must maintain and implement the operational risk management framework, including clear accountability for AI agent lifecycle (deployment, monitoring, decommissioning).	Every agent in the fleet has an owner field populated at registration. AgentGuard emits an agent.owner_assigned event and enforces that no agent can transition to 'active' without an owner. Ownership changes are logged with Owner field populated for all 18 agents.	✓ PASS	Head of AI Governance
CPS230-GOV-03	24	The entity must define and document risk appetite for operational risk, including tolerances for AI agent failures, erroneous outputs, and autonomous action limits.	AgentGuard policy files encode per-agent rate limits, cost caps, and decision thresholds. These numeric limits serve as the operationalised risk appetite. Breaches emit rate_limit.exceeded and cost_cap.exceeded events with Block+HITL rate 4.0% within tolerance.	✓ PASS	Chief Risk Officer
CPS230-GOV-04	25	The operational risk policy must explicitly cover technology and AI-related risks, including risks arising from automated decision-making and agentic AI systems.	AgentGuard's policy YAML schema enforces that each agent carries a policy_category field classifying it under the entity's risk taxonomy. Policy-category metadata is verified at evaluation time and included in every audit. 4932 evaluations all carry policy_category metadata.	✓ PASS	Head of AI Governance
CPS230-GOV-05	26	The operational risk framework must be subject to independent review, including review of controls over automated and AI-driven processes at least annually.	This evidence pack itself constitutes the artefact produced for independent review. The cryptographic Merkle root and OpenTimestamps anchor prove the event dataset was not altered post-generation. The pack is designed to Merkle root computed; OTS anchor status: pending	✓ PASS	External Auditor
CPS230-GOV-06	27	The entity must have a process to ensure lessons from operational risk incidents, including AI-related incidents, are captured and used to improve controls.	AgentGuard captures every incident.created and incident.remmediated event. The Incident and Gap Log section of this pack lists all incidents, their root cause classification, and remediation status. Repeated incident pattern Control events present in audit log.	✓ PASS	Chief Risk Officer

CPS230-GOV-07	28	The Board and senior management must foster a risk culture that supports prudent risk-taking, including staff accountability for AI agent supervision and escalation.	AgentGuard HITL escalation events capture which human approved or rejected an agent decision. The approval identity, timestamp, and decision are all logged. This demonstrates that humans are engaged in agent supervision, Rejection rate 28.0%, avg latency 253.9s — active review demonstrated.	✓ PASS	Head of AI Governance
CPS230-GOV-08	29	Material outsourcing arrangements, including those with AI platform and model providers, must be governed under a formal outsourcing framework with ongoing oversight.	AgentGuard's Third-Party and Model Registry section documents every LLM provider, model version, and API endpoint used. Model version changes emit model.version_changed events that trigger re-assessment against outsourci Control events present in audit log.	✓ PASS	Head of Vendor Risk
CPS230-OPS-01	30	The entity must identify, assess and manage material operational risks including those arising from automated processes and AI systems.	AgentGuard assigns each agent a criticality score (critical/high/medium/low) at registration. The Agent Inventory section shows the full fleet with criticality. High-criticality agents receive enhanced monitoring: all de Control events present in audit log.	✓ PASS	Head of AI Governance
CPS230-OPS-02	31	Material operational risks must be regularly assessed, with assessment covering the scope and nature of automated actions taken by AI agents.	The Tool-Call Inventory section enumerates every (agent × tool) pair with call counts. Unexpected tool combinations (agent attempting a tool outside its approved scope) generate tool.scope_violation events that are captu Control events present in audit log.	✓ PASS	Head of AI Governance
CPS230-OPS-03	32	Operational risk controls must be effective. For AI agents, this means policy controls correctly intercept out-of-policy actions at a demonstrable rate.	AgentGuard evaluates every agent action against the attached policy. The block rate, HITL escalation rate, and allow rate are computed over the review window and provided in the Statistical Sampling section. Block 2.9%, HITL 1.0% — within expected ranges.	✓ PASS	Chief Risk Officer
CPS230-OPS-04	33	The entity must have processes to identify and assess new and emerging operational risks, including risks arising from changes to AI model versions and capabilities.	Model version changes (e.g. GPT-4o-mini → GPT-4o, Claude 3 Sonnet → Claude 3.5) are captured in model.version_changed events. Each triggers a re-evaluation of the agent's risk profile before the new version is activated Control events present in audit log.	✓ PASS	Head of AI Governance
CPS230-OPS-05	34	The entity must assess and manage concentration risk in its operational risk profile, including concentration in AI model providers that could cause simultaneous failures.	The Third-Party and Model Registry section quantifies the percentage of agent evaluations routed through each provider. Concentration risk is flagged if any single provider exceeds the Board-approved concentration thresh Control events present in audit log.	✓ PASS	Head of Vendor Risk
CPS230-OPS-06	35	Material changes to operational processes, including changes to AI agent policies, must follow a formal change management process with appropriate testing and approval.	Policy version changes emit policy.version_changed events including the approver identity, prior version hash, and new version hash. AgentGuard enforces a minimum review period of 1 hour between policy upload and activat Control events present in audit log.	✓ PASS	Change Advisory Board

CPS230-OPS-07	36	The entity must have processes to detect operational risk incidents, including anomalous behaviour from AI agents that may indicate compromise or error.	AgentGuard detects anomalies using per-agent baselines. Anomalies (sudden spike in block rate, unusual tool call distribution, deviation from normal call frequency) are emitted as anomaly.detected events that trigger inc Control events present in audit log.	✓ PASS	Head of AI Governance
CPS230-OPS-08	37	The entity must maintain incident response procedures that cover AI-specific incidents including hallucination-driven actions, prompt injection attacks, and runaway automation.	Every incident in AgentGuard is classified by incident type including: policy_breach, prompt_injection_suspected, model_error, runaway_loop, and data_exfiltration_attempt. Classification drives playbook selection and SLA Control events present in audit log.	✓ PASS	Incident Response Lead
CPS230-OPS-09	38	The entity must maintain the integrity of data used in operational risk management, including the integrity of AI agent audit logs that form the basis of risk reporting.	AgentGuard implements a hash-chained audit log. Each event carries the SHA-256 hash of the previous event, forming a tamper-evident chain. The merkle.py module in this pack verifies the chain integrity across all events Control events present in audit log.	✓ PASS	External Auditor
CPS230-OPS-10	39	The entity must control costs associated with operational risk management systems. AI agent LLM spend must be within approved budgets and subject to automated rate-limiting.	Per-agent cost caps are defined in policy files. The Spend and Rate-Limit Attestation section reports total spend, per-agent cost breakdown, and any cost_cap.exceeded events in the window. Total spend: USD 29.6147. Per-agent caps active.	✓ PASS	Chief Financial Officer
CPS230-OPS-11	40	APRA must be notified of material operational risk incidents within prescribed timeframes. Material AI-related incidents (e.g. data breach via agent, regulatory non-compliance by agent) must be assessed for notification ...	Incidents classified as data_exfiltration_attempt, regulatory_breach, or customer_harm automatically generate an incident.apra_notification_candidate event flagging them for review by the CRO and Legal team. Control events present in audit log.	✓ PASS	Chief Risk Officer / Legal
CPS230-OPS-12	41	The entity must conduct root cause analysis for material operational risk incidents and implement corrective actions to prevent recurrence.	AgentGuard's incident workflow requires a root_cause field before an incident can be closed. Root cause categories include: policy_gap, model_regression, data_quality, prompt_injection, third_party_failure, and operator_ Control events present in audit log.	✓ PASS	Head of AI Governance
CPS230-OPS-13	42	For AI systems making or supporting high-risk decisions, the entity must implement human oversight controls proportionate to the risk.	AgentGuard enforces HITL gates via the hitl_required decision type in policy rules. The HITL approval log records every human decision including the reviewer identity, decision time, and whether additional context was re 50 HITL escalations; 36 approved, 14 rejected.	✓ PASS	Head of AI Governance
CPS230-OPS-14	43	The entity must continuously improve its operational risk management framework, including through periodic benchmarking of AI governance controls.	This evidence pack version is compared against the prior pack version at generation time. Changes in block rate, HITL rate, incident frequency, and coverage gaps are computed as trend indicators in the Executive Summary. Control events present in audit log.	✓ PASS	Chief Risk Officer

CPS230-BCM-01	44	The entity must maintain a business continuity plan (BCP) that covers critical business services, including AI-dependent services where agent unavailability could impair the service.	Agents classified as critical or high have a continuity_config in their registration manifest specifying fallback behaviour (fail-safe, fail-open, or manual override mode) if AgentGuard is unavailable. Continuity config Control events present in audit log.	✓ PASS	Head of Business Continuity
CPS230-BCM-02	47	The entity must identify its critical business services and the components, including AI agents and models, that support them.	Critical Business Service (CBS) mapping is maintained in agent registration metadata via the business_service field. The Agent Inventory section cross-references each agent to its CBS. Agents not mapped to a CBS are flag Control events present in audit log.	✓ PASS	Head of Business Continuity
CPS230-BCM-03	48	The entity must establish and maintain recovery time objectives (RTOs) for critical business services, including the time to restore AI agent functionality after an outage.	AgentGuard logs agent.suspended and agent.restored events. The elapsed time between these events is compared against the RTO specified in the agent's continuity_config. Control events present in audit log.	✓ PASS	Head of Business Continuity
CPS230-BCM-04	50	The entity must regularly test its BCP, including testing of AI agent failover, fallback modes, and manual override procedures.	BCP test events (continuity.test_initiated, continuity.test_completed) are logged when the ops team exercises agent failover scenarios. Test results including pass/fail and RTO achieved are recorded. Control events present in audit log.	✓ PASS	Head of Business Continuity
CPS230-BCM-05	52	The entity must identify its operational tolerances for disruptions to critical business services, including maximum tolerable outage periods for AI-supported services.	Maximum tolerable outage (MTO) is captured in agent continuity_config alongside RTO. When an agent.suspended event is detected, the system emits an alert if MTO is approached. MTO approaching events are included in the I Control events present in audit log.	✓ PASS	Head of Business Continuity
CPS230-BCM-06	55	The entity must have crisis communication procedures for significant operational disruptions, including significant AI incidents that affect customers or counterparties.	Customer-impacting AI incidents trigger an incident.customer_impact_assessed event. Where customer impact is confirmed, the incident workflow prompts issuance of a customer notification. Notification events are captured Control events present in audit log.	✓ PASS	Chief Risk Officer / Customer Relations
CPS230-SPM-01	59	The entity must identify material service providers, including AI model providers and agentic AI platform vendors, where disruption would materially affect the entity.	All AI providers are logged in the Third-Party and Model Registry. A provider is classified as material if it handles >10% of evaluations or supports any critical-agent CBS. Material provider status is reviewed quarterly Control events present in audit log.	✓ PASS	Head of Vendor Risk
CPS230-SPM-02	61	Before engaging a material service provider, the entity must conduct appropriate due diligence including assessment of security, resilience, and data handling practices.	Provider registration in AgentGuard requires attestation fields: security_posture, data_residency, soc2_certified, and penetration_test_date. These are verified at onboarding and reviewed annually. Control events present in audit log.	✓ PASS	Head of Vendor Risk

CPS230-SPM-03	62	Agreements with material service providers must include provisions addressing business continuity, data security, audit access rights, and notification of material incidents.	Provider records in AgentGuard include contract_terms_verified field, capturing whether the legal team has confirmed required contract clauses are present. Contract expiry dates are tracked and expiry warnings are generated. Control events present in audit log.	✓ PASS	Legal / Head of Vendor Risk
CPS230-SPM-04	63	The entity must monitor the ongoing performance of material service providers, including AI model providers, against agreed service levels.	AgentGuard captures provider_latency, provider_error_rate, and provider.sla_breach events. These are aggregated in the Spend and Rate-Limit Attestation section. SLA breach events trigger incident workflows. Control events present in audit log.	✓ PASS	Head of Vendor Risk
CPS230-SPM-05	65	The entity must identify and manage concentration risk arising from fourth-party dependencies, including dependencies of AI model providers on common infrastructure.	Provider registration includes infrastructure_provider field (e.g. AWS us-east-1, Azure East US). This enables detection of concentration at the infrastructure level even when multiple AI providers are used. Concentration No single region >80%. Top: us-east-1 at 67%.	✓ PASS	Head of Vendor Risk
CPS230-SPM-06	68	The entity must have documented exit strategies for material service providers, including the ability to transition AI agent workloads to an alternative provider if needed.	Each material provider record in AgentGuard includes an exit_plan_documented field and the date of the most recent exit plan test. The exit plan test is conducted annually. Control events present in audit log.	✓ PASS	Head of Vendor Risk
CPS230-SPM-07	69	Where a material service provider sub-contracts services, the entity must ensure adequate oversight of sub-contractors that are material to the service provided.	Provider records include sub_contractors list. Sub-contractor dependencies are disclosed during provider onboarding. Changes to sub-contractors trigger a provider.subcontractor_changed event requiring review. Control events present in audit log.	✓ PASS	Head of Vendor Risk
CPS230-SPM-08	71	The entity must ensure that data processed by material service providers is subject to appropriate data sovereignty controls, including controls on where AI inference occurs.	Every evaluation event carries a metadata.provider_region field. Events processed in non-approved regions generate a data_sovereignty.violation event. Approved regions are defined in the entity's data governance policy. Control events present in audit log.	✓ PASS	Chief Risk Officer / Legal

EU_AI_ACT

Control ID	Clause	Requirement	Evidence / Implementation	Result	Reviewer
EU_AI_ACT-BLOCK-PROHIBITED-MAN	—	Block subliminal manipulation techniques and exploiting vulnerabilities — Article 5 prohibited practices	AgentGuard enforces this control via the 'block' decision type for tools: subliminal_content, dark_pattern, manipulate_emotion, exploit_vulnerability_social. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-BLOCK-BIOMETRIC-MASS	—	Block real-time remote biometric identification in public spaces — Article 5(1)(d)	AgentGuard enforces this control via the 'block' decision type for tools: realtime_face_recognition, biometric_mass_scan, public_surveillance_ai, bulk_biometric_identify. 10 matching events found.	✓ PASS	Head of AI Governance

EU_AI_ACT-REQUIRE-APPROVAL-AUT	—	Require human review for consequential automated decisions about individuals — Article 14 human oversight	AgentGuard enforces this control via the 'hit' decision type for tools: credit_scoring, insurance_pricing, employment_screening, benefits_determination, risk_classification. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-REQUIRE-APPROVAL-HIG	—	Require human approval before deploying high-risk AI system outputs — Article 14	AgentGuard enforces this control via the 'hit' decision type for tools: deploy_hiring_model, deploy_credit_model, deploy_law_enforcement_ai, deploy_medical_ai. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-BLOCK-SOCIAL-SCORING	—	Block social scoring of natural persons by public authorities — Article 5(1)(c)	AgentGuard enforces this control via the 'block' decision type for tools: social_score_update, citizen_rating, social_credit_write, behaviour_score_assign. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-MONITOR-AI-GENERATED	—	Monitor AI-generated content for transparency labelling obligations — Article 50	AgentGuard enforces this control via the 'monitor' decision type for tools: generate_text, generate_image, generate_video, generate_audio, deepfake_generate. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-BLOCK-AI-TRAINING-ON	—	Block training on data lacking legal basis — Article 10 data governance	AgentGuard enforces this control via the 'block' decision type for tools: train_on_personal_data, fine_tune_without_consent, scrape_personal_profiles, harvest_biometrics. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-REQUIRE-APPROVAL-EMO	—	Require human approval for emotion recognition in workplace or education — Article 5(1)(f)	AgentGuard enforces this control via the 'hit' decision type for tools: workplace_emotion_scan, student_emotion_monitor, employee_affect_analysis. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-MONITOR-FOUNDATION-M	—	Monitor general-purpose AI model usage for systemic risk assessment — Article 55	AgentGuard enforces this control via the 'monitor' decision type for tools: call_gpt4, call_claude, call_gemini, openai_chat, anthropic_complete. 10 matching events found.	✓ PASS	Head of AI Governance
EU_AI_ACT-BLOCK-PROHIBITED-INF	—	Block inferring protected characteristics (race, political opinion, religion) from biometrics — Article 10	AgentGuard enforces this control via the 'block' decision type for tools: infer_ethnicity, infer_political_view, infer_religion_biometric, infer_sexual_orientation_ai. 10 matching events found.	✓ PASS	Head of AI Governance

ISO42001

Control ID	Clause	Requirement	Evidence / Implementation	Result	Reviewer
ISO42001-4.1-01	4.1	The organisation shall determine external and internal issues relevant to its purpose and that affect its ability to achieve the intended outcomes of its AIMS.	AgentGuard captures the deployment context of each agent at registration including: intended_use, operational_environment, and stakeholder_categories. This context is used to assess which controls apply and at what strin All 18 agents registered with required metadata.	✓ PASS	AI Governance Lead

IS042001-4.2-01	4.2	The organisation shall determine interested parties relevant to its AIMS and their requirements, including requirements related to responsible AI use and affected individuals.	Stakeholder categories (customers, regulators, employees, third parties) are captured in agent metadata. HITL decisions carry stakeholder_impact_category to record which stakeholder group a decision affects. Control events present in audit log.	✓ PASS	AI Governance Lead
IS042001-4.3-01	4.3	The organisation shall determine the boundaries and applicability of the AIMS and document the scope, including which AI systems are in scope.	The Agent Inventory section of this evidence pack constitutes the AIMS scope documentation. Every agent in scope carries the aims_in_scope=true flag in its registration metadata. All 18 agents registered with required metadata.	✓ PASS	AI Governance Lead
IS042001-5.1-01	5.1	Top management shall demonstrate leadership and commitment with respect to the AIMS, including establishing an AI policy and ensuring resources are available.	Board-approved policy versions are captured in agent registration events via the board_approved_policy_version field. The policy lineage tree is auditable. Control events present in audit log.	✓ PASS	Company Secretary
IS042001-5.2-01	5.2	The organisation shall establish an AI policy that includes commitment to responsible AI development and use, satisfying applicable requirements, and continual improvement.	AgentGuard policy files explicitly encode the responsible AI commitments: prohibited_actions, transparency_requirements, and fairness_constraints. Policy compliance is verified at every evaluation. 4932 evaluations all reference a policy_id.	✓ PASS	AI Governance Lead
IS042001-5.3-01	5.3	Top management shall ensure that the responsibilities and authorities for roles relevant to the AIMS are assigned and communicated within the organisation.	Each agent has three named roles: owner (business), reviewer (risk), and operator (technical). Role assignments are captured in agent registration and changes are tracked. All 18 agents registered with required metadata.	✓ PASS	AI Governance Lead
IS042001-6.1-01	6.1.1	The organisation shall plan, implement and maintain AI risk assessment processes, including identification of risks and opportunities associated with AI systems.	AI risk assessment is embedded in the agent registration workflow. The risk_assessment_completed event is emitted when the risk form is submitted. Risk level (critical/high/medium/low) is set based on the assessment outcome. Control events present in audit log.	✓ PASS	Head of AI Governance
IS042001-6.1-02	6.1.2	The organisation shall determine and implement appropriate risk treatments for identified AI risks, including selecting controls from Annex A.	Each risk assessment maps to specific AgentGuard policy controls (block, hitl, monitor rules). The risk_treatment_plan event records which controls were selected and why. Treatment plans are reviewable in the control map. Control events present in audit log.	✓ PASS	Head of AI Governance
IS042001-6.2-01	6.2	The organisation shall establish AI management objectives at relevant functions and levels, and plan how to achieve them.	AgentGuard captures deployment objectives in agent registration via the intended_outcomes field. Outcomes are reviewed quarterly against actual agent performance metrics. Control events present in audit log.	✓ PASS	Head of AI Governance
IS042001-7.1-01	7.1	The organisation shall determine and provide the resources needed for the establishment, implementation, maintenance and continual improvement of the AIMS.	AgentGuard's infrastructure uptime, evaluation capacity (evaluations/sec), and storage capacity are monitored. Resource constraint events (capacity.warning, capacity.critical) are captured in the audit log. Control events present in audit log.	✓ PASS	Head of Technology

ISO42001-7.2-01	7.2	The organisation shall determine and ensure that personnel doing work under the AIMS have the competence, including appropriate education, training and experience.	AgentGuard reviewer identities are cross-referenced with the HR training system via the reviewer_certified field. HITL reviewers must complete AI governance training before being allowed to approve decisions. All 50 HITL decisions made by certified reviewers.	✓ PASS	HR / Head of AI Governance
ISO42001-7.3-01	7.3	Persons doing work under the organisation's control shall be aware of the AI policy, their contribution to AIMS effectiveness, and the implications of not conforming.	Policy acknowledgement events (policy.acknowledged_by_operator) are captured when operators confirm they have read and understood the policy for agents they manage. Control events present in audit log.	✓ PASS	Head of AI Governance
ISO42001-7.4-01	7.4	The organisation shall determine the internal and external communications relevant to the AIMS, including what, when, with whom, and how to communicate.	This evidence pack is the primary internal communication artefact for the AIMS. It is generated on demand and on a scheduled basis. Pack generation events are logged. Control events present in audit log.	✓ PASS	Company Secretary
ISO42001-7.5-01	7.5	The organisation shall maintain documented information required by the AIMS, including records of AI system performance, incidents, risk assessments, and treatment decisions.	AgentGuard's immutable audit log constitutes the documented information requirement. This pack assembles the required records into a single cryptographically anchored document. Records are retained for the period specified. Control events present in audit log.	✓ PASS	External Auditor
ISO42001-8.1-01	8.1	The organisation shall plan, implement, control, monitor and review the processes needed to meet the requirements of the AIMS, and to implement the actions determined in planning.	Deployment control gates are enforced by AgentGuard: agents cannot be activated without completed risk assessment, policy attachment, owner assignment, and (for critical agents) continuity configuration. Gate compliance Control events present in audit log.	✓ PASS	Head of AI Governance
ISO42001-8.2-01	8.2	The organisation shall perform AI risk assessments at planned intervals and whenever significant changes are proposed.	Risk assessments are triggered by: new agent registration, model version changes, scope changes, and the scheduled quarterly review. The risk_assessment.triggered event captures the trigger type. Control events present in audit log.	✓ PASS	Head of AI Governance
ISO42001-8.3-01	8.3	The organisation shall implement the AI risk treatment plan and retain documented information of the risk treatment process.	Risk treatment implementation is confirmed by the policy.attached event following a risk treatment plan. The link between risk assessment, treatment plan, and policy is maintained in the audit trail. Control events present in audit log.	✓ PASS	Head of AI Governance
ISO42001-8.4-01	8.4	The organisation shall conduct impact assessments for AI systems that may have significant impacts on individuals, groups, or society.	Impact assessments are required for agents with stakeholder_categories including customers, employees, or public. The impact_assessment.completed event is required before such agents can be activated. Control events present in audit log.	✓ PASS	Head of AI Governance / Legal

ISO42001-8.5-01	8.5	The organisation shall establish and implement a process for responsible AI system development and use, including transparency in AI decision-making.	Every agent decision is logged with a reason field explaining why the decision was made. The reason is human-readable and references the specific policy rule triggered. Reason field population rate is a key transparency 4932 evaluations carry structured decision explanations.	✓ PASS	AI Governance Lead
ISO42001-9.1-01	9.1	The organisation shall evaluate the performance of its AI systems and the AIMS, and determine what needs to be monitored and measured.	AgentGuard emits daily performance reports (performance.daily_report) for each agent including: evaluation count, block rate, HITL rate, latency, error rate, and cost. These feed the Statistical Sampling section. Daily performance reports generated for 18 agents.	✓ PASS	Head of Technology
ISO42001-9.2-01	9.2	The organisation shall conduct internal audits at planned intervals to determine whether the AIMS conforms to requirements and is effectively implemented and maintained.	This evidence pack is the primary artefact for the internal AIMS audit. The Control Map section provides the audit trail for each control. The cryptographic anchoring ensures the pack is tamper-evident. Conformance 100.0%.	✓ PASS	Internal Audit / External Auditor
ISO42001-9.3-01	9.3	Top management shall review the organisation's AIMS at planned intervals to ensure its continuing suitability, adequacy, effectiveness and alignment with strategic direction.	Management review events (management_review.completed) capture when the board or executive committee has reviewed the AIMS outputs. This pack is designed to be the input document for those reviews. Control events present in audit log.	✓ PASS	Company Secretary / Board
ISO42001-10.1-01	10.1	The organisation shall continually improve the suitability, adequacy and effectiveness of the AI management system.	AgentGuard tracks improvement actions as improvement.action_created and improvement.action_completed events. These are generated from: audit findings, incident root causes, management review actions, and control gap remediation events present in audit log.	✓ PASS	Head of AI Governance
ISO42001-10.2-01	10.2	When a nonconformity occurs, the organisation shall take action to control and correct it, determine root cause, and implement corrective actions to prevent recurrence.	Nonconformities (control gaps identified in this pack) are automatically converted to improvement.action_created events with type=corrective_action. Root cause and corrective action plan are required fields before the corrective action events present in audit log.	✓ PASS	Head of AI Governance
ISO42001-A6.1-01	A.6.1	Controls shall be established for the development of AI systems, including data management, model development, testing, and deployment controls.	AgentGuard integrates with CI/CD pipelines via the agent.deployment_event that captures the source commit hash, test suite result, and approver for each agent deployment. Control events present in audit log.	✓ PASS	Head of Technology / QA
ISO42001-A6.2-01	A.6.2	The organisation shall establish controls for data used in AI systems, including data quality, provenance, and suitability for the intended AI application.	Input data quality events (data_quality.assessment_completed) capture the data provenance and quality score for datasets used by agents. Low-quality inputs generate alerts. Control events present in audit log.	✓ PASS	Head of Data Governance
ISO42001-A6.3-01	A.6.3	Where AI systems make or support consequential decisions, the organisation shall implement appropriate transparency and explainability mechanisms.	AgentGuard provides structured decision explanations for every evaluation. The explanation includes: rule_triggered, rule_description, matched_params, and confidence_score. These are visible to HITL reviewers and capture Control events present in audit log.	✓ PASS	AI Governance Lead

STATISTICAL SAMPLING

The following sample of evaluation events is drawn from the review window to demonstrate control operation. Events are selected to represent the distribution of decisions (allow / block / HITL) and tool types.

Timestamp	Agent	Tool	Params Hash	Decision	Reason (truncated)
2026-01-02 18:05:12	finance-reconcile-2	api_call	493c4e60b94c...	allow	policy \$rule:allow — tool 'api_call' is in approved scope for agent finance-reco...
2026-01-15 08:27:01	finance-reconcile-1	generate_report	80b364425d02...	allow	policy \$rule:allow — tool 'generate_report' is in approved scope for agent finan...
2026-01-20 04:57:55	ops-dispatch-1	database_query	edb915730502...	allow	policy \$rule:allow — tool 'database_query' is in approved scope for agent ops-di...
2026-01-25 15:55:31	finance-reconcile-1	send_email	37fb89dd8516...	allow	policy \$rule:allow — tool 'send_email' is in approved scope for agent finance-re...
2026-01-27 19:38:02	data-pipeline-1	database_query	02a4832bf360...	allow	policy \$rule:allow — tool 'database_query' is in approved scope for agent data-p...
2026-01-29 11:06:45	customer-support-2	http_post	eac148302364...	allow	policy \$rule:allow — tool 'http_post' is in approved scope for agent customer-su...
2026-01-30 05:41:23	finance-reconcile-2	send_email	5ffe24e52333...	allow	policy \$rule:allow — tool 'send_email' is in approved scope for agent finance-re...
2026-02-09 13:54:11	data-pipeline-1	database_query	594288e6c27b...	allow	policy \$rule:allow — tool 'database_query' is in approved scope for agent data-p...

2026-02-10 22:57:35	ops-dispatch-2	read_file	53e1f03ee36b...	allow	policy \$rule:allow — tool 'read_file' is in approved scope for agent ops-dispatc...
2026-02-14 19:09:14	ops-dispatch-3	read_file	4aec6099284b...	allow	policy \$rule:allow — tool 'read_file' is in approved scope for agent ops-dispatc...
2026-02-23 09:41:23	customer-support-2	database_query	b8bd3fcc6fe6...	allow	policy \$rule:allow — tool 'database_query' is in approved scope for agent custom...
2026-02-25 06:59:16	finance-reconcile-1	generate_report	49e492ceffa3...	block	policy \$rule:block-unauthorised-action — tool call outside approved scope
2026-02-25 22:06:00	sales-research-2	send_email	1ee32868e635...	allow	policy \$rule:allow — tool 'send_email' is in approved scope for agent sales-rese...
2026-03-04 15:51:49	research-analyst-1	read_file	fb6d06eef0c6...	allow	policy \$rule:allow — tool 'read_file' is in approved scope for agent research-an...
2026-03-07 09:32:57	research-analyst-1	database_query	3826e5d82425...	allow	policy \$rule:allow — tool 'database_query' is in approved scope for agent resear...
2026-03-08 00:46:06	brand-content-2	read_file	c0d173f18f58...	allow	policy \$rule:allow — tool 'read_file' is in approved scope for agent brand-conte...
2026-03-21 12:00:22	ops-dispatch-1	send_email	4fecd92599cf...	allow	policy \$rule:allow — tool 'send_email' is in approved scope for agent ops-dispat...
2026-03-24 15:28:18	sales-research-1	web_search	e8e2b09cef9e...	allow	policy \$rule:allow — tool 'web_search' is in approved scope for agent sales-rese...

2026-03-25 21:15:56	customer-support-2	read_file	fa40018d14a2...	allow	policy \$rule:allow — tool 'read_file' is in approved scope for agent customer-su...
2026-03-26 08:46:13	infra-monitor-2	api_call	bc02532876ea...	allow	policy \$rule:allow — tool 'api_call' is in approved scope for agent infra-monito...

Sample: 20 events drawn from 4,932 total evaluations (0.41% sampling rate).

INCIDENT AND GAP LOG

All block/denied evaluations generating incidents and all HITL escalations in the window are listed below, with incident type, severity, and remediation status.

Timestamp	Agent	Type	Severity	Status	Root Cause / Note
2026-01-05 10:04:47	brand-content-1	policy breach	MEDIUM	RESOLVED	model_regression
2026-01-06 07:11:32	customer-support-3	policy breach	MEDIUM	RESOLVED	data_quality
2026-01-10 05:02:02	infra-monitor-2	prompt injection suspected	HIGH	RESOLVED	third_party_failure
2026-01-10 22:44:15	sales-research-1	prompt injection suspected	HIGH	RESOLVED	third_party_failure
2026-01-13 07:10:33	brand-content-1	tool scope violation	MEDIUM	OPEN	Under investigation
2026-01-15 17:54:41	finance-reconcile-2	tool scope violation	MEDIUM	RESOLVED	policy_gap
2026-01-17 02:09:19	customer-support-2	policy breach	MEDIUM	RESOLVED	third_party_failure
2026-01-21 04:06:08	infra-monitor-2	tool scope violation	MEDIUM	OPEN	Under investigation
2026-01-25 12:20:38	data-pipeline-1	data exfiltration attempt	HIGH	RESOLVED	model_regression
2026-01-27 20:47:25	brand-content-1	data exfiltration attempt	HIGH	RESOLVED	data_quality
2026-01-29 15:31:19	customer-support-1	model error	MEDIUM	OPEN	Under investigation
2026-02-02 18:16:47	finance-reconcile-1	tool scope violation	MEDIUM	RESOLVED	policy_gap
2026-02-07 07:17:47	customer-support-3	data exfiltration attempt	HIGH	RESOLVED	model_regression
2026-02-13 22:40:07	data-pipeline-1	policy breach	MEDIUM	RESOLVED	data_quality
2026-02-15 05:56:04	brand-content-2	model error	MEDIUM	RESOLVED	third_party_failure
2026-02-16 19:47:10	research-analyst-1	tool scope violation	MEDIUM	RESOLVED	policy_gap
2026-02-17 20:46:28	ops-dispatch-1	prompt injection suspected	HIGH	RESOLVED	third_party_failure
2026-02-18 00:33:52	ops-dispatch-2	data exfiltration attempt	HIGH	RESOLVED	model_regression
2026-02-19 11:20:28	infra-monitor-2	tool scope violation	MEDIUM	RESOLVED	model_regression
2026-02-21 21:47:41	infra-monitor-2	tool scope violation	MEDIUM	RESOLVED	policy_gap
2026-02-28 07:21:49	research-analyst-1	data exfiltration attempt	HIGH	RESOLVED	model_regression
2026-03-09 21:49:17	brand-content-1	tool scope violation	MEDIUM	RESOLVED	model_regression
2026-03-11 15:59:10	customer-support-2	tool scope violation	MEDIUM	RESOLVED	policy_gap
2026-03-13 18:38:36	data-pipeline-1	policy breach	MEDIUM	RESOLVED	data_quality

2026-03-15 01:06:44	finance-reconcile-1	model error	MEDIUM	RESOLVED	operator_error
2026-03-15 17:26:36	sales-research-1	prompt injection suspected	HIGH	RESOLVED	operator_error
2026-03-15 23:36:50	ops-dispatch-1	data exfiltration attempt	HIGH	RESOLVED	policy_gap
2026-03-16 13:26:54	finance-reconcile-2	policy breach	MEDIUM	RESOLVED	third_party_failure
2026-03-20 03:32:12	customer-support-3	prompt injection suspected	HIGH	RESOLVED	operator_error
2026-03-21 17:38:29	customer-support-1	policy breach	MEDIUM	RESOLVED	operator_error
2026-03-22 11:13:26	sales-research-1	prompt injection suspected	HIGH	RESOLVED	data_quality
2026-03-23 15:32:50	sales-research-2	data exfiltration attempt	HIGH	RESOLVED	policy_gap
2026-03-24 08:08:26	sales-research-1	tool scope violation	MEDIUM	RESOLVED	third_party_failure
2026-03-30 06:26:39	ops-dispatch-2	tool scope violation	MEDIUM	OPEN	Under investigation
2026-03-30 17:48:34	finance-reconcile-1	data exfiltration attempt	HIGH	OPEN	Under investigation

Total: 35 incidents. Resolved: 30. Open: 5.

THIRD-PARTY AND MODEL REGISTRY

All LLM providers, model versions, and inference regions used during the review window. Data residency notes reflect inference location, not storage location.

Provider	Model Version	Inference Region	Jurisdiction	Agents	Data Residency Note
Anthropic	claude-3-5-haiku-20241022	Australia (Sydney)	Australia	2	Privacy Act 1988 §8 applies. Data processed in-country.
Anthropic	claude-3-5-sonnet-20241022	Australia (Sydney)	Australia	1	Privacy Act 1988 §8 applies. Data processed in-country.
Anthropic	claude-3-5-sonnet-20241022	USA (Virginia)	United States	1	GDPR: Standard Contractual Clauses apply. Privacy Act 1988: cross-border disclosure consent obtained.
Anthropic	claude-3-haiku-20240307	Australia (Sydney)	Australia	2	Privacy Act 1988 §8 applies. Data processed in-country.
Google	gemini-1.5-flash	USA (Iowa)	United States	1	GDPR: Standard Contractual Clauses apply.
Openai	gpt-4o	USA (Virginia)	United States	7	GDPR: Standard Contractual Clauses apply. Privacy Act 1988: cross-border disclosure consent obtained.
Openai	gpt-4o-mini	USA (Virginia)	United States	4	GDPR: Standard Contractual Clauses apply. Privacy Act 1988: cross-border disclosure consent obtained.

SPEND AND RATE-LIMIT ATTESTATION

Total evaluations (window)	4,932
Total AI provider spend (USD)	\$29.6147
Average cost per evaluation	\$0.006005
Rate-limit ceiling (fleet)	1,000 evaluations/agent/day (configurable per-agent)
Cost cap per agent (daily)	USD 10.00 (configurable per-agent)
Cost cap breach events	0
Rate-limit breach events	0

Evaluation distribution by provider:

Provider	Evaluations	Share	Est. Cost (USD)
Openai	3,014	61.1%	\$9.04
Anthropic	1,644	33.3%	\$4.11
Google	274	5.6%	\$0.22

TOOL-CALL INVENTORY AND APPROVAL LOG

All unique (agent × tool) pairs observed in the window with call counts. This section supports tool governance reviews and assists in identifying unexpected tool usage.

Tool	Call Count	% of Total
database_query	960	19.5%
read_file	831	16.8%
web_search	537	10.9%
send_email	517	10.5%
http_post	480	9.7%
api_call	467	9.5%
write_file	292	5.9%
http_get	242	4.9%
generate_text	145	2.9%
generate_report	105	2.1%
monitor_metrics	97	2.0%
crm_update	81	1.6%
initiate_payment	58	1.2%
send_alert	54	1.1%
data_transform	45	0.9%
create_image	21	0.4%

SIGNATURE AND ANCHOR PROOF

This section provides the cryptographic proofs that authenticate this evidence pack. The Merkle root is computed over all events included in the pack. The OpenTimestamps receipt anchors this root to the Bitcoin blockchain, providing an independent, immutable timestamp that cannot be backdated.

Cryptographic Summary

Source audit-log SHA-256	853aa49c9ff4e449840e9d90075f21620ad53b7a307c01f0afd476c0e429aafe
Merkle root (all events)	642ad10aefb062d4c54d54287d1aef2827cdd5c98854be49fa111afe9c647cf9
PDF SHA-256	computed post-generation
Hash chain integrity	pass (5041 events, 0 broken links)
OTS anchor status	pending
OTS calendar	—
OTS file	—
Generated at	2026-05-15 13:47 UTC
Generator version	agentguard-evidence v1.0.0

BTC Anchor Status: PENDING — BTC anchor pending — OTS error: No module named 'opentimestamps.serialize'. Retry with: agentguard-evidence anchor

To complete anchoring after network connectivity is restored, run: agentguard-evidence anchor evidence-pack.pdf

Verification Instructions

- To verify the integrity of this evidence pack independently:
1. Install: `pip install agentguard-evidence opentimestamps-client`
 2. Run: `agentguard-evidence verify evidence-pack.pdf`
 3. This will: (a) recompute and compare the Merkle root, (b) verify the hash chain, (c) verify the OTS receipt against a Bitcoin block header, (d) confirm the PDF SHA-256.
 4. A passing verification confirms the dataset was not modified after the pack was generated.